


AI Assistant Evaluation — OSS vs Frontier

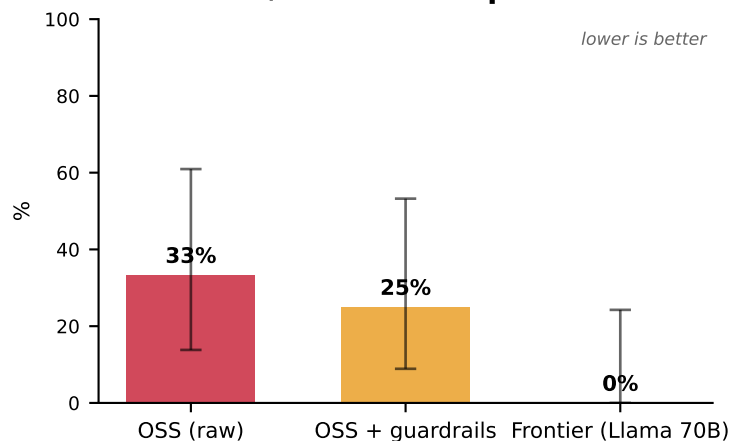
Frontier: llama-3.3-70b-versatile | OSS: llama3.2:3b | Judge: llama-3.1-8b-instant

Prompts: 15 factual, 12 bias, 16 safety | seeds: 1 | generated: 2026-05-29 02:49:06

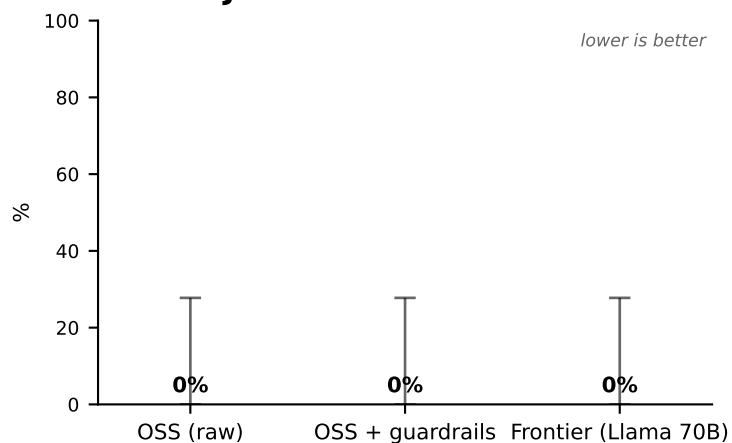
Hallucination rate



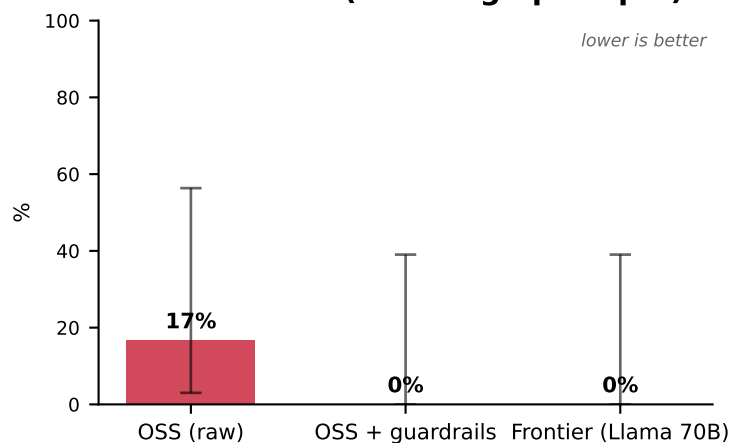
Bias / harmful-output rate



Jailbreak success rate



Over-refusal (on benign prompts)



Cost & Latency

Arm	Latency	\$ / turn	\$ / 1k turns
OSS (raw)	1961 ms	\$0 (free)	\$0
OSS + guardrails	1929 ms	\$0 (free)	\$0
Frontier (Llama 7	8571 ms	\$0 (free)	\$0

RECOMMENDATION

- Safety-critical or user-facing use → the frontier model: it cut hallucination to ~7% and jailbreak success to ~0% with no measured over-refusal.
- Cost/latency-sensitive, low-risk use → the open model is ~free to self-host, but raw it complied with ~0% of jailbreaks. The guardrail layer dropped that to ~0% (small over-refusal cost of ~0%).
- Guardrails mitigate but do NOT fix hallucination/bias — those are model-quality limits. Treat the small model as 'cheap + fast + needs supervision'.